

# Lab 5: Parallel training of neural networks

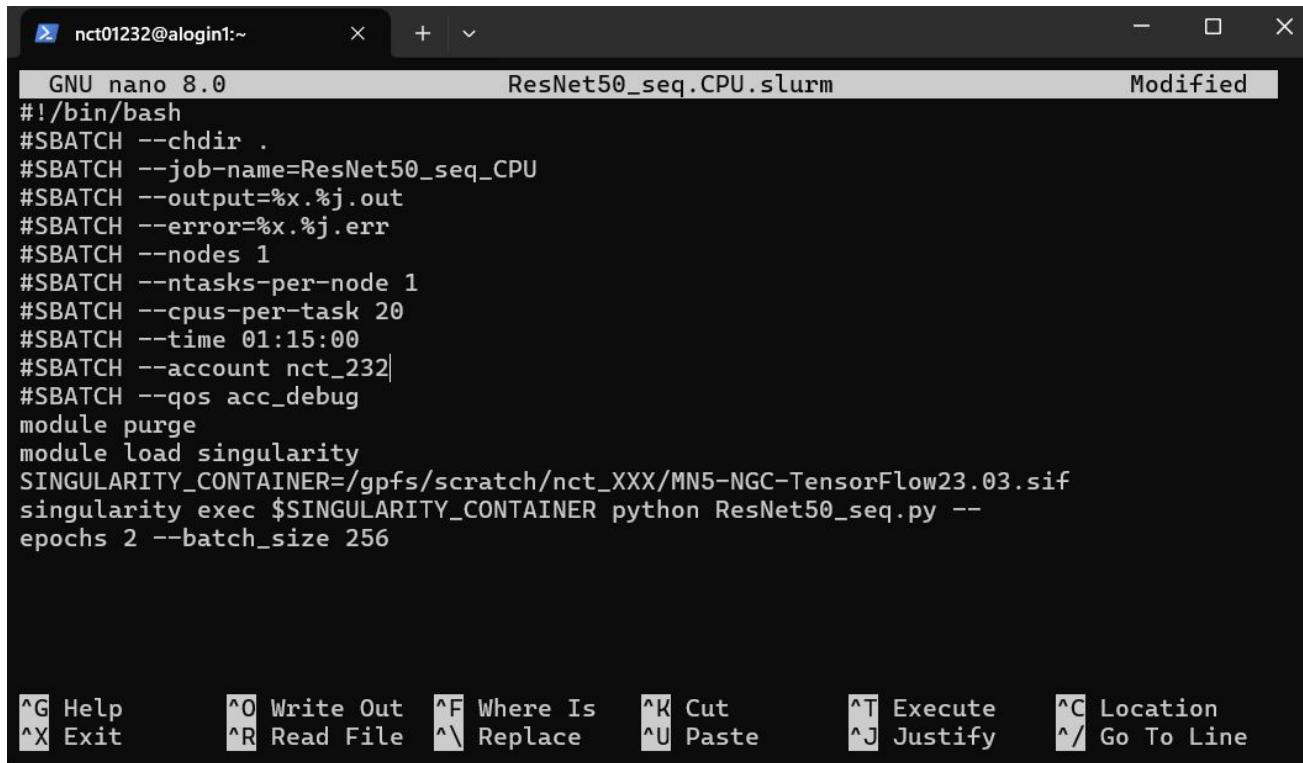


Mario Ventura

# 1. Code Overview

- **Argumentos:** 5 épocas, 2048 tamaño lote
- **Datos:** CIFAR-10 redimensionado en 128x128
- **Modelo:** ResNet 50V2 para 10 clases
- **Entrenamiento:** Usa optimizador SGD

## 2. Entrenamiento en CPU



The image shows a terminal window with a dark background. The title bar at the top indicates the user is 'nct01232@alodin1' in their home directory. The terminal is running GNU nano 8.0, editing a file named 'ResNet50\_seq.CPU.slurm'. The script content is as follows:

```
#!/bin/bash
#SBATCH --chdir .
#SBATCH --job-name=ResNet50_seq_CPU
#SBATCH --output=%x.%j.out
#SBATCH --error=%x.%j.err
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 1
#SBATCH --cpus-per-task 20
#SBATCH --time 01:15:00
#SBATCH --account nct_232
#SBATCH --qos acc_debug
module purge
module load singularity
SINGULARITY_CONTAINER=/gpfs/scratch/nct_XXX/MN5-NGC-TensorFlow23.03.sif
singularity exec $SINGULARITY_CONTAINER python ResNet50_seq.py --
epochs 2 --batch_size 256
```

At the bottom of the terminal, there is a help bar with various keyboard shortcuts for nano editor functions:

<b>^G</b> Help	<b>^O</b> Write Out	<b>^F</b> Where Is	<b>^K</b> Cut	<b>^J</b> Execute	<b>^C</b> Location
<b>^X</b> Exit	<b>^R</b> Read File	<b>^_</b> Replace	<b>^U</b> Paste	<b>^I</b> Justify	<b>^_</b> Go To Line

## 2. Entrenamiento en CPU

```
nct01232@alogin1:~  
[nct01232@alogin1 ~]$ sbatch ResNet50_seq.CPU.slurm  
Submitted batch job 17116453  
[nct01232@alogin1 ~]$ squeue --start  
      JOBID PARTITION    NAME    USER  ST       START_TIME   NODES SCHEDNODES          NODELIST(REASON)  
      17116453      acc ResNet50  nct01232  PD               N/A         1 (null)             (Priority)  
[nct01232@alogin1 ~]$ |
```

```
[nct01232@alogin1 ~]$ ls -l  
total 5  
-rw-r--r-- 1 nct01232 nct 1538 Mar  9 18:45 MNIST.py  
-rw-r--r-- 1 nct01232 nct  450 Mar 10 17:57 MNIST.slurm  
-rw-r--r-- 1 nct01232 nct  467 Mar 17 18:10 ResNet50_seq.CPU.slurm  
-rw-r--r-- 1 nct01232 nct  407 Mar 17 18:10 ResNet50_seq_CPU.17116453.err  
-rw-r--r-- 1 nct01232 nct    0 Mar 17 18:10 ResNet50_seq_CPU.17116453.out  
drwxr-xr-x 2 nct01232 nct 4096 Mar  8 14:10 lab4  
drwxr-xr-x 2 nct01232 nct 4096 Mar  9 19:07 singularity_cache  
drwxr-xr-x 2 nct01232 nct 4096 Mar  9 19:07 singularity_tmp  
[nct01232@alogin1 ~]$ |
```

## 2. Entrenamiento en CPU

Diferencia: 9s

```
[nct01232@alogin1 ~]$ cat ResNet50_seq_CPU.17117579.out  
Epoch 1/2  
196/196 - 246s - loss: 2.0349 - accuracy: 0.2521 - 246s/epoch - 1s/step  
Epoch 2/2  
196/196 - 241s - loss: 1.7586 - accuracy: 0.3552 - 241s/epoch - 1s/step  
[nct01232@alogin1 ~]$ |
```

## 2. Entrenamiento en CPU

Diferencia: 9s

```
[nct01232@alogin1 ~]$ cat ResNet50_seq_CPU.17117579.out
Epoch 1/2
196/196 - 246s - loss: 2.0349 - accuracy: 0.2521 - 246s/epoch - 1s/step
Epoch 2/2
196/196 - 241s - loss: 1.7586 - accuracy: 0.3552 - 241s/epoch - 1s/step
[nct01232@alogin1 ~]$ |
```

- Entrenar un modelo sin GPUs no es práctico
- Puede ser que la etapa (epoch) 1 tarde más por la carga inicial



nct01232@alagin1:~

Windows PowerShell

GNU nano 8.0

ResNet50\_seq.CPU.slurm

Modified

```
#!/bin/bash
#SBATCH --chdir .
#SBATCH --job-name=ResNet50_seq_GPU
#SBATCH --output=%x.%j.out
#SBATCH --error=%x.%j.err
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 1
#SBATCH --cpus-per-task 20 #Minimum cpus requested should
# be (nodes * gpus/node * 20)
#SBATCH --time 01:15:00
#SBATCH --account nct_325
#SBATCH --qos acc_debug
#SBATCH --gres gpu:1

echo " "
echo "Job Name: $SLURM_JOB_NAME"
echo "Job ID: $SLURM_JOB_ID"
echo "Submit Directory: $SLURM_SUBMIT_DIR"
echo "Partition: $SLURM_JOB_PARTITION"
echo "Account: $SLURM_JOB_ACCOUNT"
echo " "
echo "Number of Nodes Allocated: $SLURM_JOB_NUM_NODES"
echo "Total number of tasks: $SLURM_NTASKS"
echo "Number of tasks per node: $SLURM_NTASKS_PER_NODE"
echo "Number of CPU cores per task: $SLURM_CPUS_PER_TASK"
echo " "
echo "Specific GPUs allocated:"

nvidia-smi
module purge
module load singularity/4.1.5

SINGULARITY_CONTAINER=/gpfs/scratch/nct_325/MN5-NGC-TensorFlow-23.03.sif
singularity exec --nv $SINGULARITY_CONTAINER python ResNet50_seq.py --epochs 1 --batch_size 256
```

^G Help

^X Exit

^O Write Out

^R Read File

^F Where Is

^\_ Replace

^K Cut

^U Paste

^T Execute

^J Justify

^C Location

^\_ Go To Line

M-U Undo

M-E Redo

M-A Set Mark

M-6 Copy

M-] To Bracket

^B Where Was



## 2. Entrenamiento en GPU

Specific GPUs allocated:

Tue Mar 18 16:29:19 2025

NVIDIA-SMI 535.86.10			Driver Version: 535.86.10			CUDA Version: 12.2		
GPU	Name	Perf	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp		Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute	M. MIG M.
0	NVIDIA H100		On	00000000:2C:00.0	Off		0	
N/A	38C	P0	65W / 700W	2MiB / 65247MiB		0%	Default	Disabled

Processes:								
GPU	GI	CI	PID	Type	Process name		GPU Memory	
	ID	ID					Usage	
No running processes found								

196/196 - 24s - loss: 2.0328 - accuracy: 0.2547 - 24s/epoch - 125ms/step

[nct01232@alodin1 ~]\$

# 3. CPU vs GPU

**CPU:** 246s / época

**GPU:** 24s / época

- 11x más rápido en GPU.
- 2083 imágenes/s (GPU) vs 172 imágenes/s (CPU)

### 3. CPU vs GPU

**CPU:** 246s / época

**GPU:** 24s / época

- 11x más rápido en GPU.
- 2083 imágenes/s (GPU) vs 172 imágenes/s (CPU)
- **Conclusión: GPU is the winner**

# 4. Métricas de rendimiento

## 4.1 ACELERACIÓN (speedrun)

Razón entre tiempo secuencial y paralelo (1 vs n GPUs)

- Ejemplo: “1 GPU takes 80 seconds, and with 4 GPUs it takes 25 seconds, then: Speedup=80/25=3.2x”
- Ideal: **4x**

# 4. Métricas de rendimiento

## 4.2 THROUGHPUT

Velocidad con la que “algo” se procesa.

Velocidad **EFFECTIVA**.

- Proporciona **información valiosa para evaluar rendimiento computacional**.
- Ejemplo: imágenes procesadas / t.
- $100,000 \text{ imágenes} / 100\text{s} = 1000 \text{ imgs/s}$ .

# 4. Métricas de rendimiento

## 4.3 ESCALABILIDAD

La capacidad de un sistema para manejar eficientemente cantidades de trabajo crecientes/decrecientes.

Capacidad de escalar (manejar más trabajo).

Depende mayoritariamente de factores como configuraciones, tipo de arquitectura de la red, eficiencia del framework de deep learning, etc.

## 4. Métricas de rendimiento

Son de gran utilidad para evaluar la eficiencia, escalabilidad, etc. de entrenamientos de modelos con multi-GPU

# 5. Entrenamiento paralelo (TF)

## Mirrored Strategy

- Permite aplicar la estrategia llamada “Mirrored Strategy”
- tf.distribute.MirroredStrategy permite entrenamiento paralelo
- *tf.distribute.MirroredStrategy* está incluido en el conjunto de estrategias de tf.distribute.Strategy

```
mirrored_strategy = tf.distribute.MirroredStrategy()
```



# 5. Entrenamiento paralelo (TF)

## Funcionamiento del Mirrored Strategy

- Replicar el modelo en cada GPU
- Cada GPU procesa un subconjunto de los datos.
- Actualizaciones sincronizadas

```
devices = tf.config.experimental.list_physical_devices("GPU")
```

```
mirrored_strategy =  
tf.distribute.MirroredStrategy(devices=["/gpu:0", "/gpu:1"])
```

# 6. Ejecución paralela de ResNet50

```
[nct01232@alagin1 ~]$ cat ResNet50.slurm
#!/bin/bash
#SBATCH --chdir .
#SBATCH --job-name=ResNet50
#SBATCH --output=%x.%j.out
#SBATCH --error=%x.%j.err
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 1
#SBATCH --gres gpu:4
#SBATCH --cpus-per-task 80
#SBATCH --time 01:15:00
#SBATCH --account nct_325
#SBATCH --qos acc_debug

echo "PARALLEL EXECUTION USING 1, 2, 4 GPUS"
echo " "
echo "Job Name: $SLURM_JOB_NAME"
echo "Job ID: $SLURM_JOB_ID"
echo "Submit Directory: $SLURM_SUBMIT_DIR"
echo "Partition: $SLURM_JOB_PARTITION"
echo "Account: $SLURM_JOB_ACCOUNT"
echo " "

echo "Number of Nodes Allocated: $SLURM_JOB_NUM_NODES"
echo "Total number of tasks: $SLURM_NTASKS"
echo "Number of tasks per node: $SLURM_NTASKS_PER_NODE"
echo "Number of CPU cores per task: $SLURM_CPUS_PER_TASK"
echo " "
echo "Specific GPUs allocated:"
nvidia-smi

module purge
module load singularity/4.1.5

SINGULARITY_CONTAINER=/gpfs/scratch/nct_325/MN5-NGC-TensorFlow-23.03.sif

singularity exec --nv $SINGULARITY_CONTAINER python ResNet50.py --epochs 5 --batch_size 2048 --n_gpus 1
singularity exec --nv $SINGULARITY_CONTAINER python ResNet50.py --epochs 5 --batch_size 4096 --n_gpus 2
singularity exec --nv $SINGULARITY_CONTAINER python ResNet50.py --epochs 5 --batch_size 8192 --n_gpus 4
[nct01232@alagin1 ~]$
```

# 6. Ejecución paralela de ResNet50

```
[nct01232@allogin1 ~]$ sbatch ResNet50.slurm
```

```
Submitted batch job 17187550
```

```
[nct01232@allogin1 ~]$ squeue --start
```

JOBID	PARTITION	NAME	USER	ST	START_TIME	NODES	SCHEDNODES	NODELIST(REASON)
17187550	acc	ResNet50	nct01232	PD	N/A	1	(null)	(None)

```
[nct01232@allogin1 ~]$ |
```

```
[nct01232@allogin1 ~]$ ls -l
```

```
total 7
```

```
-rw-r--r-- 1 nct01232 nct 1538 Mar  9 18:45 MNIST.py
-rw-r--r-- 1 nct01232 nct  450 Mar 10 17:57 MNIST.slurm
-rw-r--r-- 1 nct01232 nct  624 Mar 18 17:42 ResNet50.17187550.err
-rw-r--r-- 1 nct01232 nct 3116 Mar 18 17:42 ResNet50.17187550.out
-rw-r--r-- 1 nct01232 nct 1202 Mar 18 17:35 ResNet50.py
-rw-r--r-- 1 nct01232 nct 1219 Mar 18 17:39 ResNet50.slurm
-rw-r--r-- 1 nct01232 nct 1026 Mar 18 16:28 ResNet50_seq.CPU.slurm
-rw-r--r-- 1 nct01232 nct  848 Mar 17 18:44 ResNet50_seq.py
drwxr-xr-x 2 nct01232 nct 4096 Mar  8 14:10 lab4
drwxr-xr-x 2 nct01232 nct 4096 Mar  9 19:07 singularity_cache
drwxr-xr-x 2 nct01232 nct 4096 Mar  9 19:07 singularity_tmp
[nct01232@allogin1 ~]$ |
```

# 6. Ejecución paralela de ResNet50

```
Epoch 1/5
25/25 - 29s - loss: 2.2495 - accuracy: 0.1618 - 29s/epoch - 1s/step
Epoch 2/5
25/25 - 14s - loss: 2.1402 - accuracy: 0.2211 - 14s/epoch - 544ms/step
Epoch 3/5
25/25 - 14s - loss: 2.0531 - accuracy: 0.2533 - 14s/epoch - 541ms/step
Epoch 4/5
25/25 - 14s - loss: 1.9943 - accuracy: 0.2773 - 14s/epoch - 545ms/step
Epoch 5/5
25/25 - 13s - loss: 1.9144 - accuracy: 0.3053 - 13s/epoch - 540ms/step
Epoch 1/5
13/13 - 32s - loss: 2.2787 - accuracy: 0.1514 - 32s/epoch - 2s/step
Epoch 2/5
13/13 - 8s - loss: 2.1381 - accuracy: 0.2145 - 8s/epoch - 623ms/step
Epoch 3/5
13/13 - 7s - loss: 2.0521 - accuracy: 0.2506 - 7s/epoch - 575ms/step
Epoch 4/5
13/13 - 8s - loss: 1.9814 - accuracy: 0.2790 - 8s/epoch - 623ms/step
Epoch 5/5
13/13 - 8s - loss: 1.9227 - accuracy: 0.3049 - 8s/epoch - 613ms/step
Epoch 1/5
7/7 - 47s - loss: 2.3040 - accuracy: 0.1436 - 47s/epoch - 7s/step
Epoch 2/5
7/7 - 5s - loss: 2.1782 - accuracy: 0.1977 - 5s/epoch - 729ms/step
Epoch 3/5
7/7 - 5s - loss: 2.0979 - accuracy: 0.2302 - 5s/epoch - 703ms/step
Epoch 4/5
7/7 - 4s - loss: 2.0452 - accuracy: 0.2549 - 4s/epoch - 632ms/step
Epoch 5/5
[nct01232@alagin1 ~]$ |
```

# 6. Ejecución paralela de ResNet50

**1 GPU**

**Promedio:** 13'75s/época

**Throughput:**  $50.000 / 13'75 \approx 3.636 \text{ img/s}$

# 6. Ejecución paralela de ResNet50

2 GPUs

Promedio: 7'75s/época

Throughput:  $50.000 / 7'75 \approx 6.452 \text{ img/s}$

# 6. Ejecución paralela de ResNet50

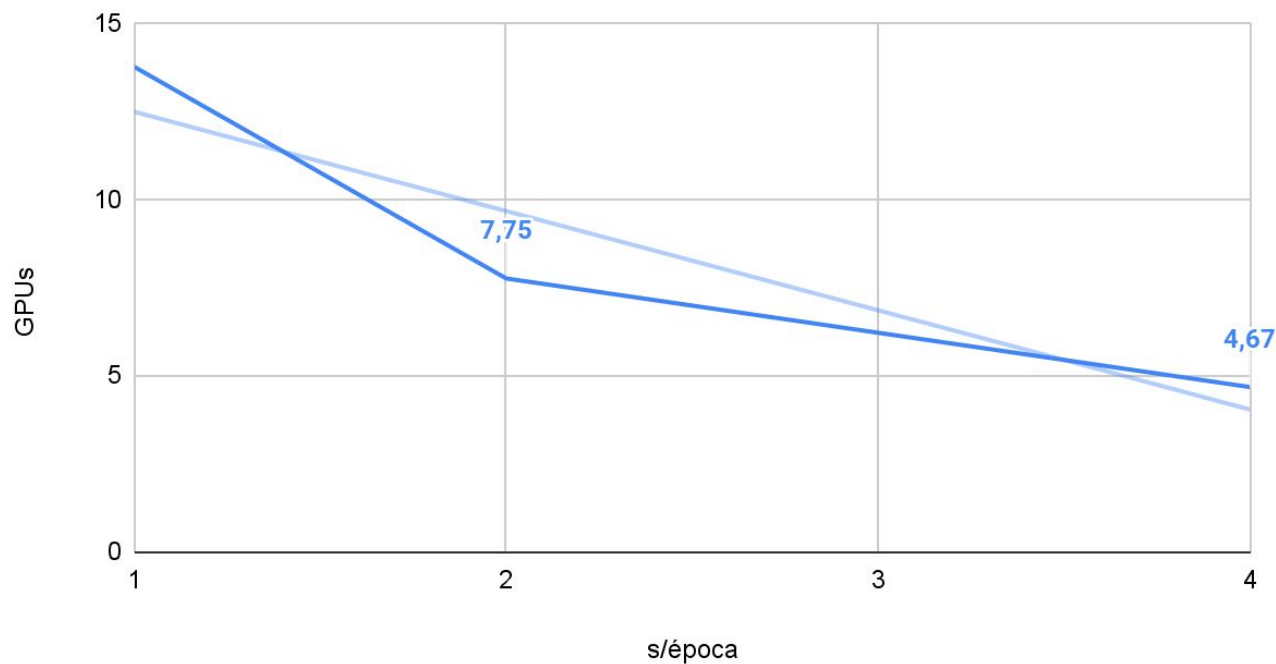
4 GPUs

Promedio: 4'67s/época

Throughput:  $50.000 / 4'67 \approx 10.707$  img/s

# 6. Ejecución paralela de ResNet50

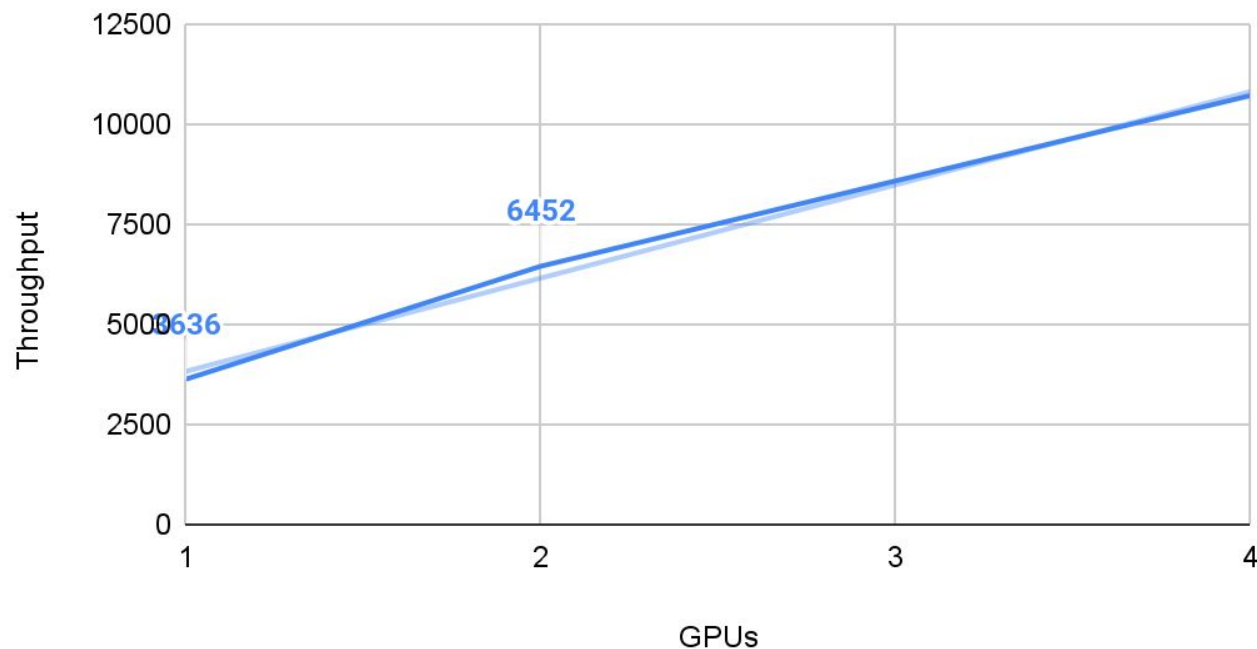
GPUs frente a s/época





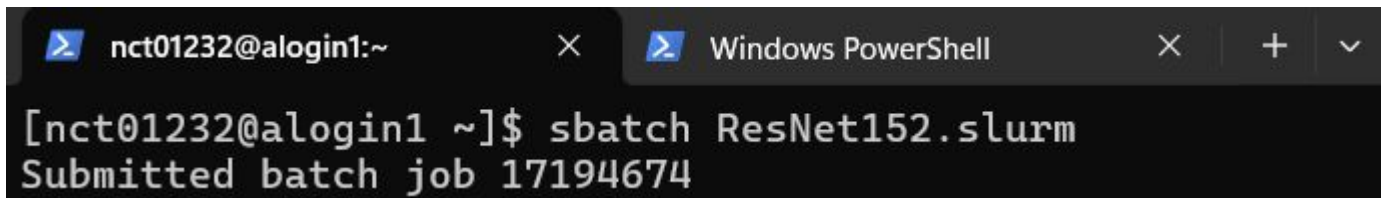
# 6. Ejecución paralela de ResNet50

Throughput frente a GPUs



# 7. ResNet152

- Modifier `.slurm`
- Modifier `.py`
- Disminuir batch size (512, 1024, 2048)



A terminal window with two tabs: 'nct01232@alogin1:~' and 'Windows PowerShell'. The active tab shows the command `sbatch ResNet152.slurm` being executed, resulting in the output 'Submitted batch job 17194674'.

```
nct01232@alogin1:~$ sbatch ResNet152.slurm
Submitted batch job 17194674
```

# 7. ResNet152

Información de archivos .err


```
key: "replicate_on_split"
value {
  b: false
}
}
experimental_type {
  type_id: TFT_PRODUCT
  args {
    type_id: TFT_DATASET
    args {
      type_id: TFT_PRODUCT
      args {
        type_id: TFT_TENSOR
        args {
          type_id: TFT_UINT8
        }
      }
    }
    args {
      type_id: TFT_TENSOR
      args {
        type_id: TFT_UINT8
      }
    }
  }
}
}
}

2025-03-18 19:11:50.345170: I tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:428] Loaded cuDNN version 8801
2025-03-18 19:11:51.588974: I tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:648] TensorFlow-32 will be used for the matrix multiplication. This
will only be logged once.
2025-03-18 19:11:51.881842: I tensorflow/compiler/xla/service/service.cc:173] XLA service 0xba0df70 initialized for platform CUDA (this does not guarantee t
hat XLA will be used). Devices:
2025-03-18 19:11:51.881881: I tensorflow/compiler/xla/service/service.cc:181] StreamExecutor device (0): NVIDIA H100, Compute Capability 9.0
2025-03-18 19:11:51.883501: I tensorflow/compiler/xla/service/service.cc:181] StreamExecutor device (1): NVIDIA H100, Compute Capability 9.0
2025-03-18 19:11:51.883508: I tensorflow/compiler/xla/service/service.cc:181] StreamExecutor device (2): NVIDIA H100, Compute Capability 9.0
2025-03-18 19:11:51.883511: I tensorflow/compiler/xla/service/service.cc:181] StreamExecutor device (3): NVIDIA H100, Compute Capability 9.0
2025-03-18 19:11:52.181796: I tensorflow/compiler/jit/xla_compilation_cache.cc:480] Compiled cluster using XLA! This line is logged at most once for the li
fetime of the process.
[nct01232@alogn1 ~]$
```

```
Epoch 1/5
98/98 - 62s - loss: 2.2160 - accuracy: 0.1742 - 62s/epoch - 637ms/step
Epoch 2/5
98/98 - 34s - loss: 2.0167 - accuracy: 0.2574 - 34s/epoch - 342ms/step
Epoch 3/5
98/98 - 33s - loss: 1.8923 - accuracy: 0.3058 - 33s/epoch - 341ms/step
Epoch 4/5
98/98 - 34s - loss: 1.7938 - accuracy: 0.3440 - 34s/epoch - 342ms/step
Epoch 5/5
98/98 - 33s - loss: 1.7049 - accuracy: 0.3751 - 33s/epoch - 341ms/step
Epoch 1/5
49/49 - 75s - loss: 2.2119 - accuracy: 0.1739 - 75s/epoch - 2s/step
Epoch 2/5
49/49 - 19s - loss: 2.0039 - accuracy: 0.2596 - 19s/epoch - 379ms/step
Epoch 3/5
49/49 - 18s - loss: 1.9089 - accuracy: 0.2987 - 18s/epoch - 371ms/step
Epoch 4/5
49/49 - 18s - loss: 1.8232 - accuracy: 0.3368 - 18s/epoch - 369ms/step
Epoch 5/5
49/49 - 18s - loss: 1.7447 - accuracy: 0.3650 - 18s/epoch - 372ms/step
Epoch 1/5
25/25 - 124s - loss: 2.2577 - accuracy: 0.1657 - 124s/epoch - 5s/step
Epoch 2/5
25/25 - 10s - loss: 2.0728 - accuracy: 0.2417 - 10s/epoch - 399ms/step
Epoch 3/5
25/25 - 10s - loss: 1.9614 - accuracy: 0.2797 - 10s/epoch - 394ms/step
Epoch 4/5
25/25 - 10s - loss: 1.8309 - accuracy: 0.3267 - 10s/epoch - 400ms/step
Epoch 5/5
25/25 - 10s - loss: 1.7452 - accuracy: 0.3537 - 10s/epoch - 403ms/step
[nct01232@alogin1 ~]$ |
```


# 7. ResNet152

s/época frente a GPU



# 7. ResNet152

Throughput frente a GPU



# 8. ResNet152 vs ResNet50

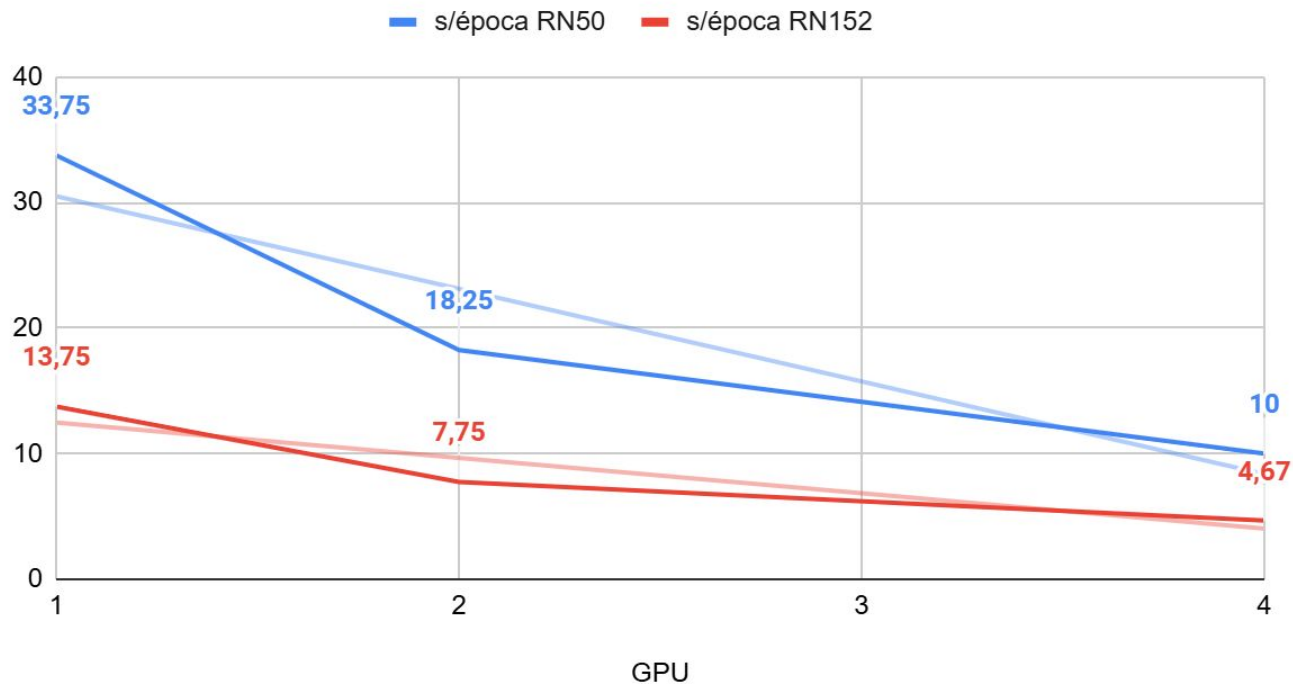
Comparación de resultados

ResNet152V2		
GPU	s/época	Throughput
1	33,75	1481
2	18,25	2739
4	10	5000

ResNet50		
GPUs	s/época	Throughput
1	13,75	3636
2	7,75	6452
4	4,67	10.707

# 8. ResNet152 vs ResNet50

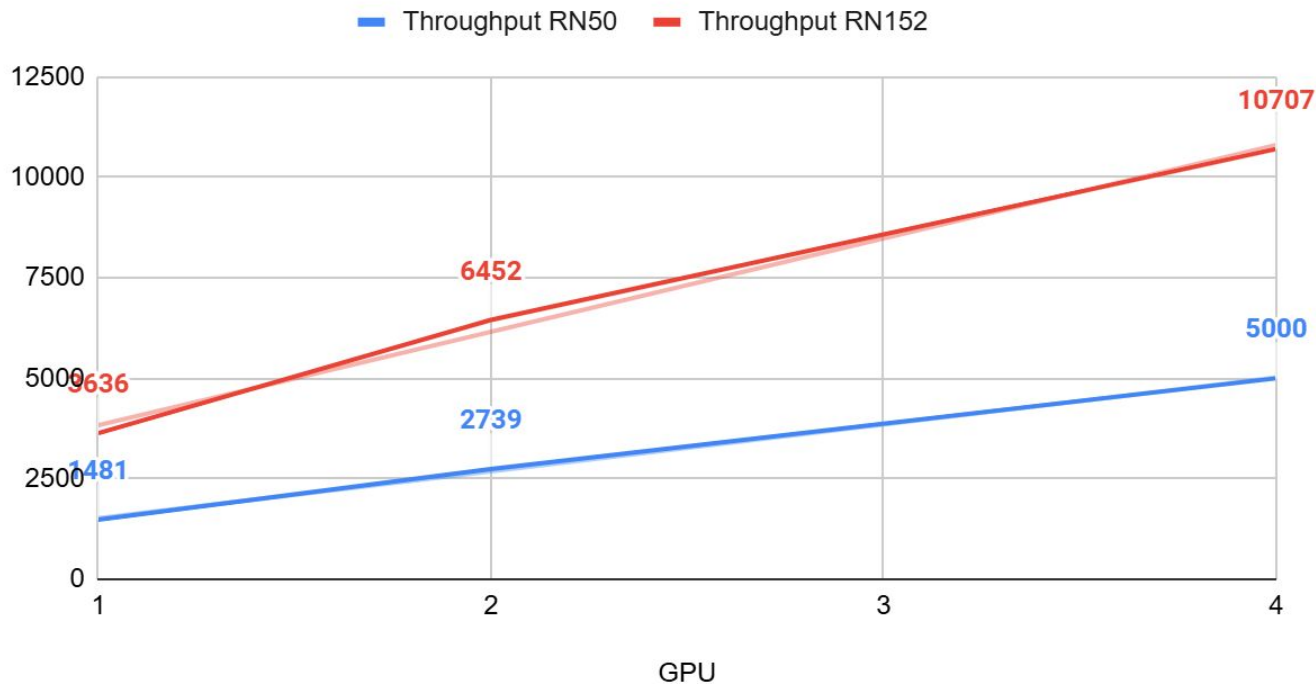
Diferencias en segundos por época





# 8. ResNet152 vs ResNet50

Diferencias en Throughput



**Gracias**